

PV&PC

Friday, January 7, 2022 10:30 AM

How to use Persistent Volume and Persistent Claims | Kubernetes

Working with kubernetes is always fun as well as challenging. The more you dive deep into the kubernetes ecosystem the more you learn.

It always bugged me when I started working with kubernetes that - How can I retain the data after the end of pod life cycle?

Answer is -

Kubernetes Persistent Volume and Persistent claims help you to retain the data of the pod even after the end of the pod life cycle

What problems does it solve?

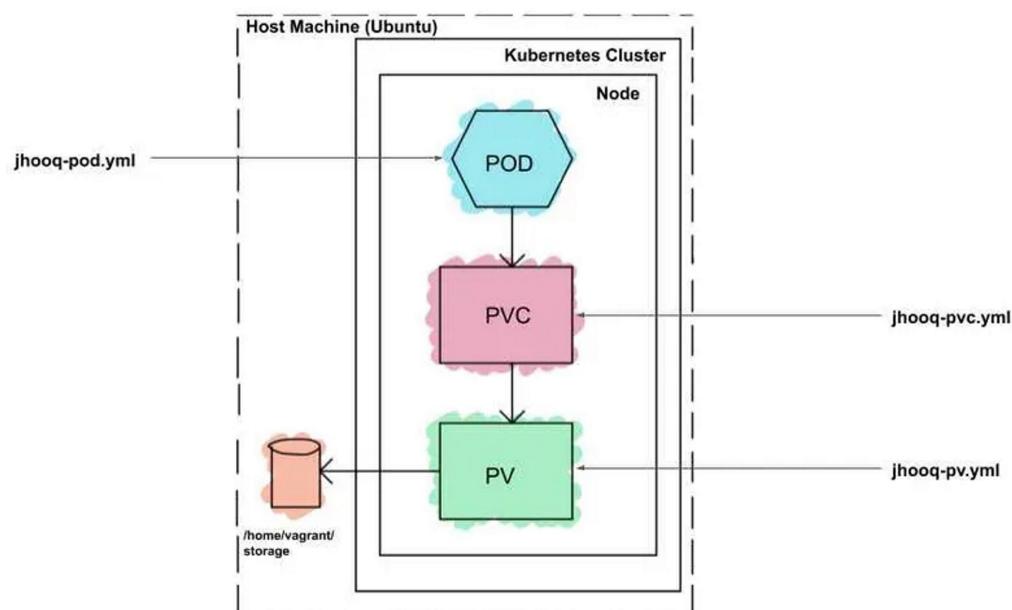
Containers running inside the pod can not share the files with each other.

1. All the files inside the container are temporary which means if you terminate the container you are going to lose all your files.
2. Secondly if in any case, your container crashes then there is no way to recover files.

Kubernetes provides volume plugin as **Persistent Volume** to address the above problems.

The lifecycle of these volumes are independent of the lifecycle of pods.

So if PODs are terminated then volumes are unmounted and detached keeping the data intact.



What is Persistent Volume(PV)?

In simple terms, it's storage available within your Kubernetes cluster. This storage can be provisioned by you or Kubernetes administrator.

It's basically a directory with some data in it and all the containers running inside the pods can access it. But Persistent Volumes are independent of the POD life cycle.

So if PODs live or die, persistent volume does get affected and it can be reused by some other PODs.

Kubernetes provides many volume plugins based on the cloud service provider you are using -
[awsElasticBlockStore](#), [azureDisk](#), [azureFile](#), [cephfs](#), [cinder](#), [configMap](#), [csi](#), [downwardAPI](#), [emptyDir](#), [fc](#) (fibre channel), [flexVolume](#), [flocker](#), [gcePersistentDisk](#), [gitRepo](#) (deprecated), [glusterfs](#), [hostPath](#), [iscsi](#), [local](#), [nfs](#), [persistentVolumeClaim](#), [projected](#), [portworxVolume](#), [quobyte](#), [rbd](#), [scaleIO](#), [secret](#), [storageos](#), [vsphereVolume](#)

How can you Create persistent volume?

There are some prerequisites before you create your persistent volume

Step 1- Prerequisites

- Kubernetes Cluster:** - You should have Kubernetes cluster up and running (**If you do not know "How to setup then please refer to [this article](#)**)
- Docker Image/container:** - A docker image that can be deployed as Kubernetes deployment. (I am going to use Spring Boot Docker image. If you do not have then please refer on **how to create docker image you can follow [this article](#)**)
- Directory for persistent Volume storage:-** Create one directory .i.e. **/home/vagrant/storage** inside your Linux machine for persistent volume. You can keep the directory name and path as per your need

Step 2 - Create a persistent volume configuration (**jhooq-pv.yml**)

Its fairly easy to create it. Refer to the following configuration which you can customize with your requirements -

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: jhooq-demo-pv
spec:
  capacity:
    storage: 1Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  storageClassName: local-storage
  local:
    path: /home/vagrant/storage
  nodeAffinity:
    required:
      nodeSelectorTerms:
        - matchExpressions:
            - key: kubernetes.io/hostname
              operator: In
              values:
                - node1
```

I saved the above configuration with name **jhooq-pv.yml** but you can assign any name of your choice.

In step 1 we have created Persistent Volume for Local Storage, now you need to apply the configuration.

Step 3 - Apply the configuration

```
kubectl apply -f jhooq-pv.yml
```

BASH

```
persistentvolume/jhooq-demo-pv created
```

BASH

Step 4 - Lets check the status of PV ↴

BASH

```
kubectl get pv
```

BASH

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON
jhooq-demo-pv	1Gi	RWO	Retain	Available		local-storage	

Here is the screen of the status

Great now you have created PV, let's move ahead and create a Persistent Volume claim.

What is Persistent Volume claim?

Persistent volume provides you an abstraction between the consumption of the storage and implementation of the storage.

In the nutshell you can say its a request for storage on behalf of an application which is running on cluster.

How to use Persistent Volume claim(PVC) ?

If you as an application developer wants to use/access Persistent Volume(PV) then you must create a request for storage and it can be done by creating PVC objects.

Step 1 - Alright lets create your first Persistent Volume Claim(jhooc-pvc.yml) -

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: jhooc-pvc
spec:
  volumeName: jhooc-demo-pv
  storageClassName: local-storage
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

YAML

Save the above configuration with file name of your choice, in mycase I am saving this file with the name **jhooc-pvc.yml**

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: jhooc-pvc
spec:
  volumeName: jhooc-demo-pv
  storageClassName: local-storage
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

Step 2 - Now apply this configuration using the following command

```
kubectl create -f jhooq-pvc.yml
```

BASH

```
persistentvolumeclaim/jhooq-pvc created
```

BASH

Step 3 - Lets check our Persistent Volume and persistent Volume Claim

```
kubectl get pvc
```

BASH

```
vagrant@kmaster:~$ kubectl get pvc
NAME      STATUS    VOLUME          CAPACITY   ACCESS MODES  STORAGECLASS     AGE
jhooq-pvc  Bound     jhooq-demo-pv  1Gi        RWO          local-storage  17h
```

BASH

```
vagrant@kmaster:~$ kubectl get pvc
NAME      STATUS    VOLUME          CAPACITY   ACCESS MODES  STORAGECLASS     AGE
jhooq-pvc  Bound     jhooq-demo-pv  1Gi        RWO          local-storage  17h
```

Create a POD using Persistent Volume claim

Now in this step we are going to create a POD using the PV and PVC from the previous steps.

We are going to deploy Spring Boot Docker image but you can use any docker application of your choice. But if you do not have any docker image with you then you can refer to - [How to deploy spring boot application in Kubernetes cluster](#)

Step 1 - Create POD configuration yml .i.e. - “jhooq-pod.yml” ↴

```
apiVersion: v1
kind: Pod
metadata:
  name: jhooq-pod-with-pvc
  labels:
    name: jhooq-pod-with-pvc
spec:
  containers:
    - name: jhooq-pod-with-pvc
      image: rahulwagh17/kubernetes:jhooq-k8s-springboot
      ports:
        - containerPort: 8080
          name: www
      volumeMounts:
        - name: www-persistent-storage
          mountPath: /home/vagrant/storage
    volumes:
      - name: www-persistent-storage
        persistentVolumeClaim:
          claimName: jhooq-pvc
```

```
apiVersion: v1
kind: Pod
metadata:
  name: jhooq-pod-with-pvc
  labels:
    name: jhooq-pod-with-pvc
spec:
  containers:
    - name: jhooq-pod-with-pvc
      image: rahulwagh17/kubernetes:jhooq-k8s-springboot
      ports:
        - containerPort: 8080
          name: www
      volumeMounts:
        - name: www-persistent-storage
          mountPath: /home/vagrant/storage
    volumes:
      - name: www-persistent-storage
        persistentVolumeClaim:
          claimName: jhooq-pvc
```

Step 2 : Lets apply the pod configuration

```
kubectl apply -f jhooq-pod.yml
```

BASH

```
pod/jhooq-pod-with-pvc created
```

BASH

```
vagrant@kmaster:~/jhooq-demo-pv$ kubectl apply -f jhooq-pod.yml  
pod/jhooq-pod-with-pvc created
```

Lets check the pod status

```
$ kubectl get pod
```

BASH

NAME	READY	STATUS	RESTARTS	AGE
jhooq-pod-with-pvc	1/1	Running	0	8m37s

BASH

Now you have deployed your pod successfully using the Persistent Volume and Persistent Volume Claim.

Step 3 : Test the microservice deployed under the POD

In this testing step we need to access the microservice which we deployed inside the POD.

To test the POD first we need to find the IP address on which it is running.

Use the following kubectl command to find the IP address of the POD

```
$ kubectl get pod -o wide
```

BASH

It should return you with (You may get different IP address) -

```
vagrant@node1:~$ kubectl get pod -o wide  
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES  
jhooq-pod-with-pvc  1/1     Running   0          96m    10.233.90.3  node1  <none>        <none>
```

Ref : <https://jhoog.com/>